


RAPPORT PROJET TP (Calcule Parallel)

1 - Parallel Scheme :



- The diagram represents k nodes, one master (Node #0), the others slaves
- The master node owns the entire database as well as the instances/centroids
- The slaves receive the data (for the database only a part) from the master, they perform the KNN/K-Means calculations and send the results back to master for results display
- The slaves can have possible communication between them, example : the case of MPI_Allreduce

Information of my machine :

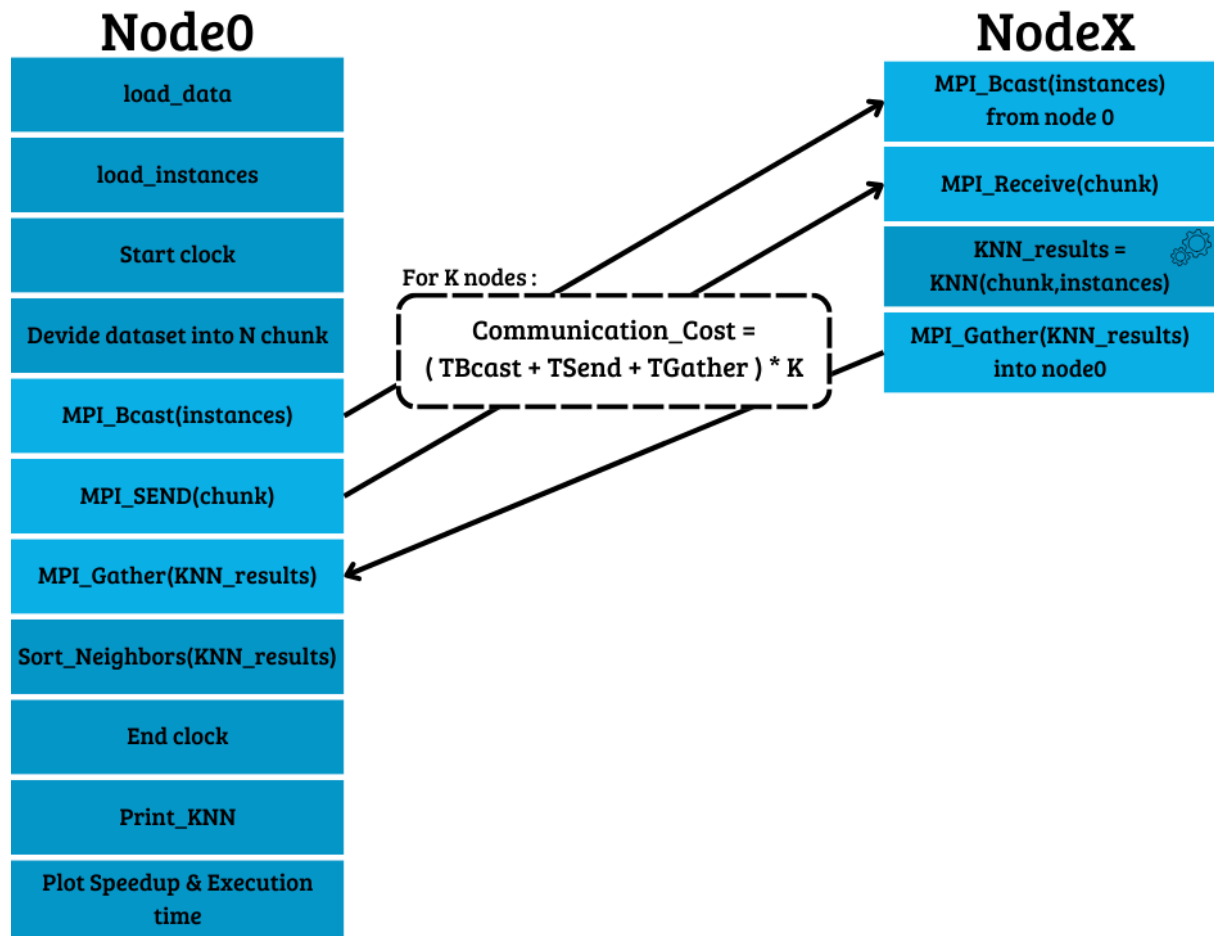
CPU : Processor Intel(R) Core(TM) i7-8650U CPU @ 1.90GHz, 2112 Mhz, 4 Core(s), 8 Logical Processor(s)

Memory : Installed Physical Memory (RAM) 16.0 GB

2 - MPI :

- MPI \Rightarrow Distributed memory
- Communication cost due to data exchange between processes
- When process 0 sends data, the process receiver will create a copy of that data in his own memory space, which it can modify independently without affecting the data in other processes.

KNN :



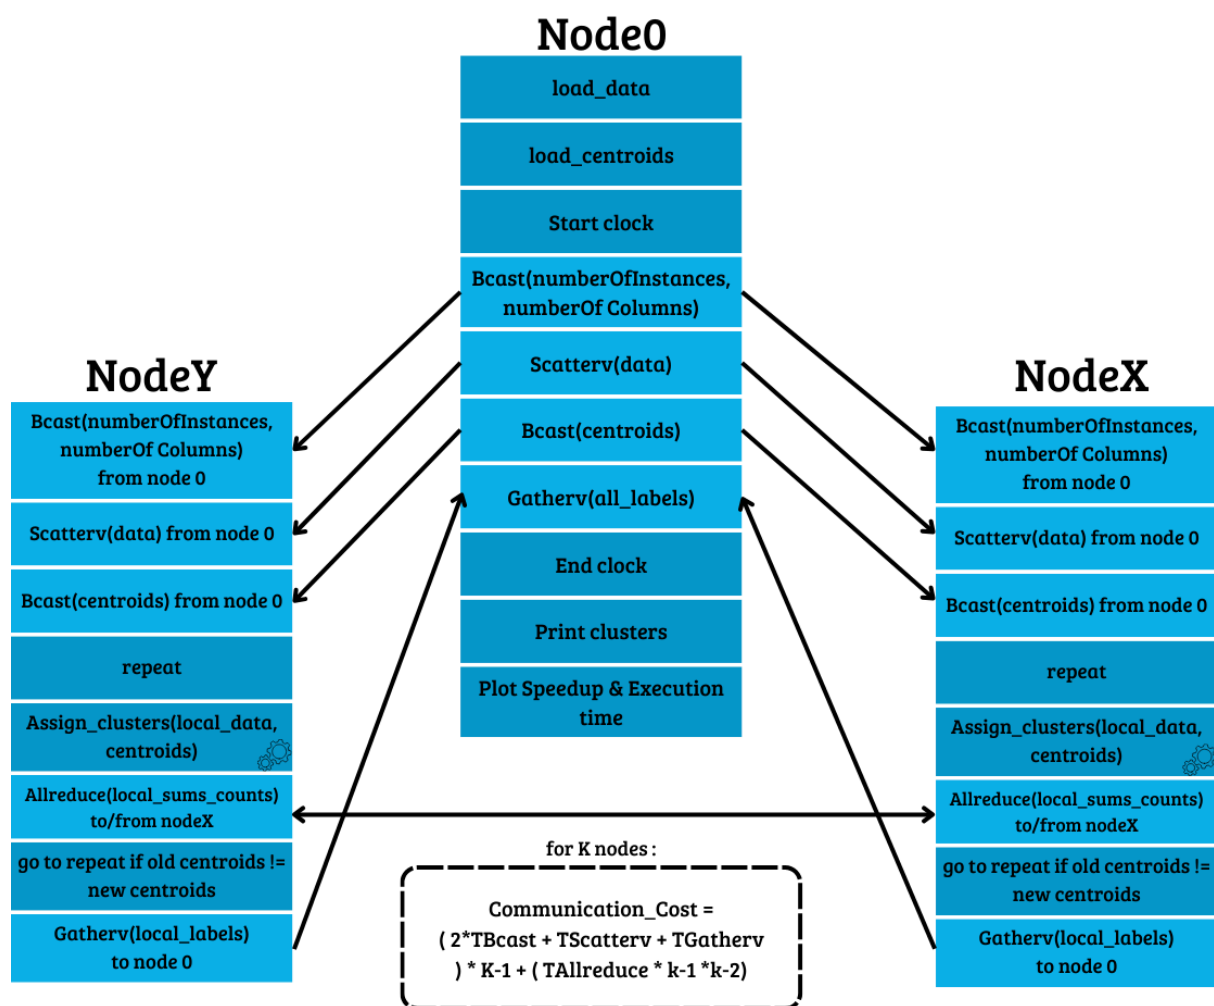
MPI_Bcast	MPI_Send	MPI_Receive	MPI_Gather
Is used to distribute a message (data) from one process (node 0) to all other processes, each receiver will have a copy of the message in his own memory space \Rightarrow collective communication	Is used to send a message from one process to another \Rightarrow point-to-point communication	Is used to receive a message sent by another process	Is used to collect data from all processes and gather it into a single process.

Variables description :

Let's consider a cluster of N nodes :

	Description	Size
Instances	Instances to classify	$W*1$ = nbr of attributes (in my case I classify one instance at a time)
chunk	part of the dataset	$L/N*(W+1)$ L : nbr of instances in the dataset W : nbr of attributes + class label
knn_results	each process writes the k nearest neighbors that he found	$K*size$ size = nbr of processes

K-Means :



MPI_Bcast	MPI_Scatterv	MPI_Allreduce	MPI_Gatherv
Same	Is used to send data from the root process to all processes. It's useful when sending data with varying lengths.	Is useful when performing a reduction operation (ex SUM) across all processes and distributing the result back to all processes.	Same as gather + data with potentially different lengths

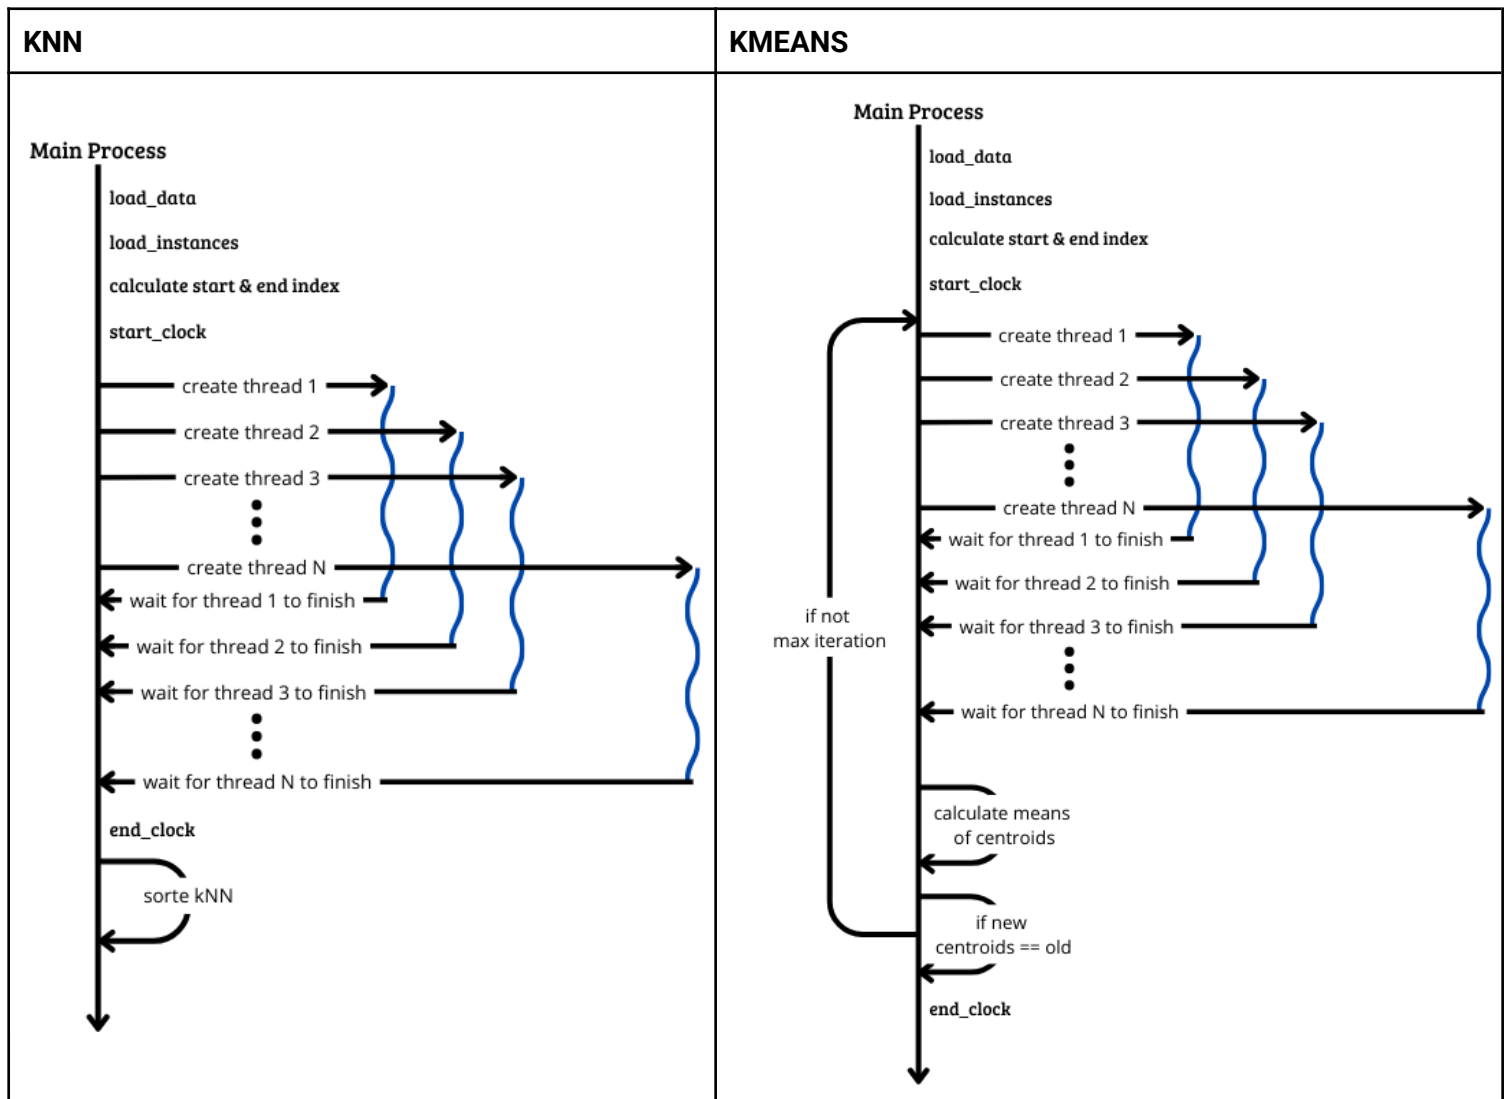
Variables description :

Let's consider a cluster of N nodes :

	Description	Size
data	Dataset	$L \times W$ L : nbr of instances W : nbr of attributes
centroids	Gravity centers (C)	$C \times (W+1)$ <i>+1 because of class label</i>
local_data	Part of data	$\sim L/N \times W$
local_sums_counts	local sums : Each entry in the matrix, <code>local_sums[c][w]</code> , represents the sum of the <i>w</i> -th attribute of all data points currently assigned to the <i>c</i> -th cluster.	$C \times W$
	local counts : how many instances are in each cluster	C
local_labels	The <i>i</i> th entry represent the class of the <i>i</i> th instance	L/N
all_labels	local_labels all gathered	L

3 - Threads :

- Shared memory
- No communication cost




Data passed for each thread :

KNN	K-means
Dataset	Dataset
Starting index	Starting index
Ending index	Ending index
Instances to classify	Centroids
Array to write KNN results (write k nearest neighbors only)	Array to write k-means results (labels)

- Each thread executes knn/ k-means and write results in knn/k-means results

4 - Curves :

I've worked with "[covertime](#)" dataset :



Description :

Classification into 7 forest cover types based on attributes such as elevation, aspect, slope, hillshade, soil-type, and more.

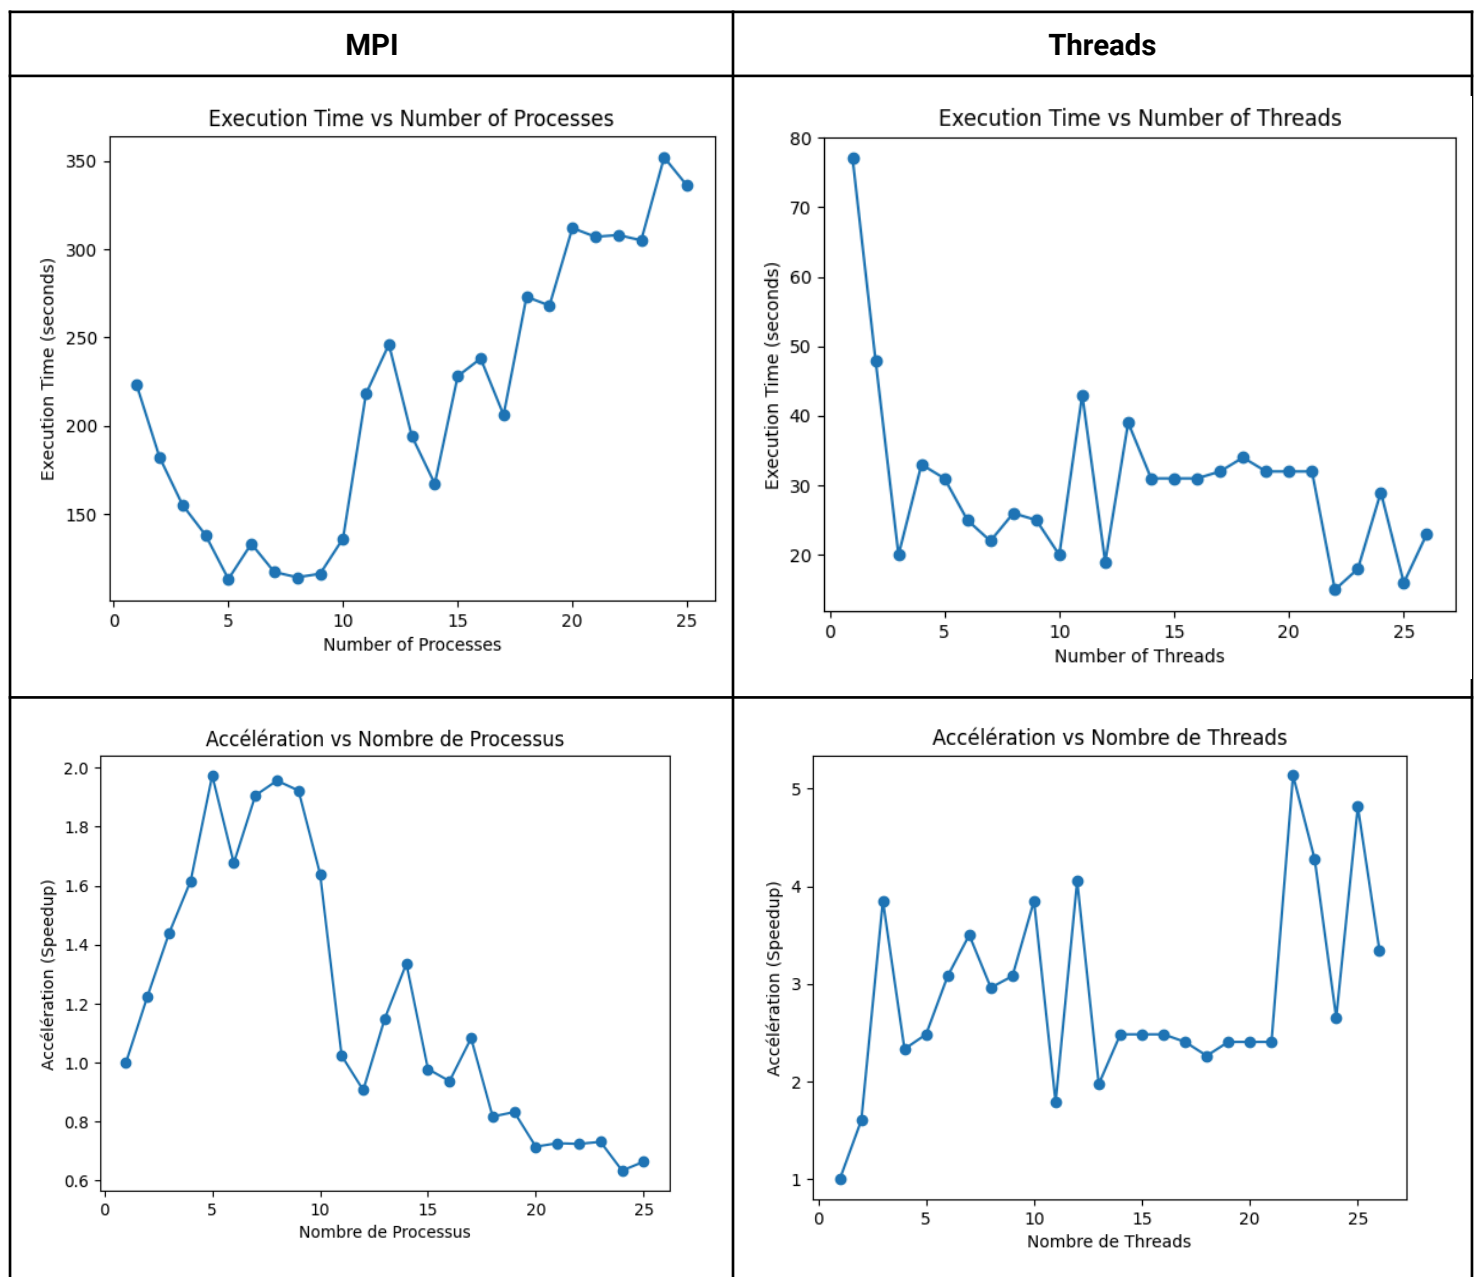
Number of Instances : 581011 (+58k)

Number of Attributes : 54 + 1 class label

7 types of class labels :

Spruce/Fir, Lodgepole Pine, Ponderosa Pine, Cottonwood/Willow, Aspen, Douglas-fir, Krummholz

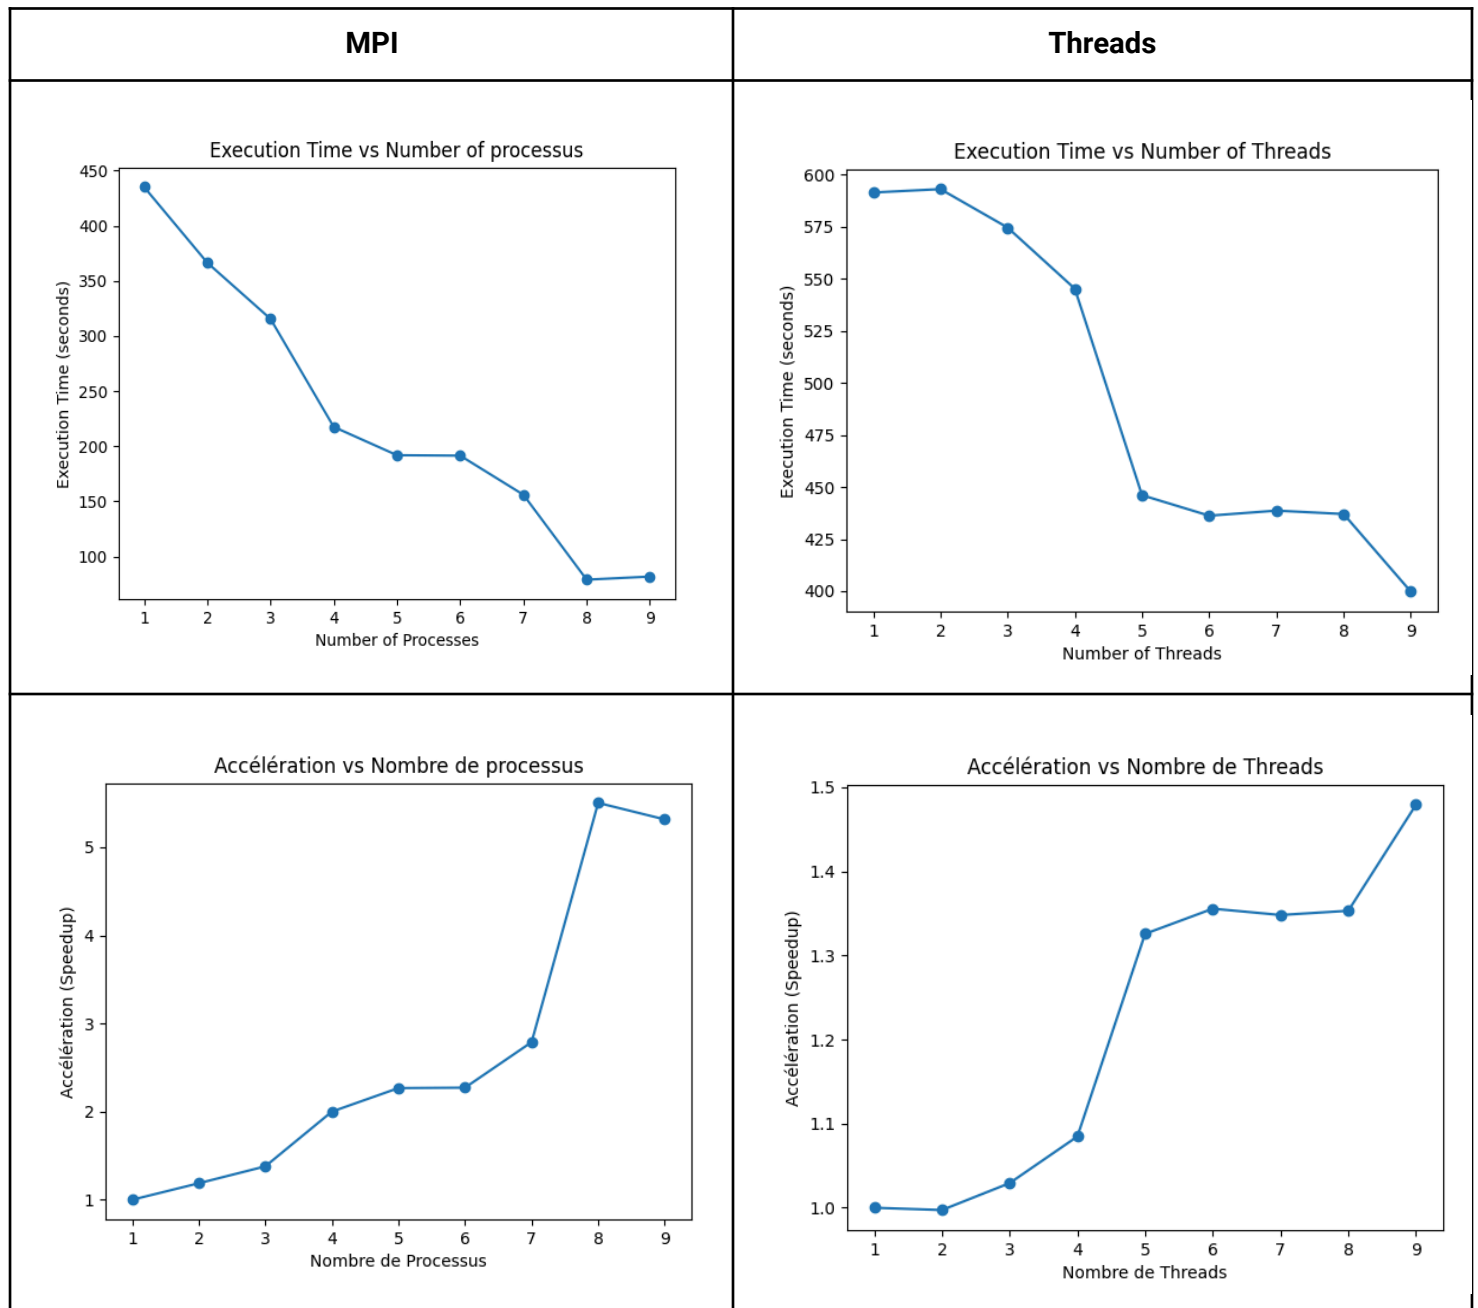
KNN :



	Peak
MPI	5 processes
THREADS	22 processes

Threads vs Processes : Threads show better results than MPI in term of execution time and speedup

K-Means :



Threads vs Processes : MPI show better results than THREADS in term of execution time and speedup